

7 Dumb Ways to Bypass Syscall Filters

Brian Welch



About me

- In between developer gigs
- Find my projects at <https://brianwel.ch>

Disclaimer

- Only talking about Linux
 - I know nothing about syscall filtering on Windows
- Mainly focusing on the amd64 architecture
- Mainly focusing on `seccomp-bpf` as the syscall filter mechanism
- Most of my testing was done on Ubuntu 20.04 (Linux kernel 5.4)

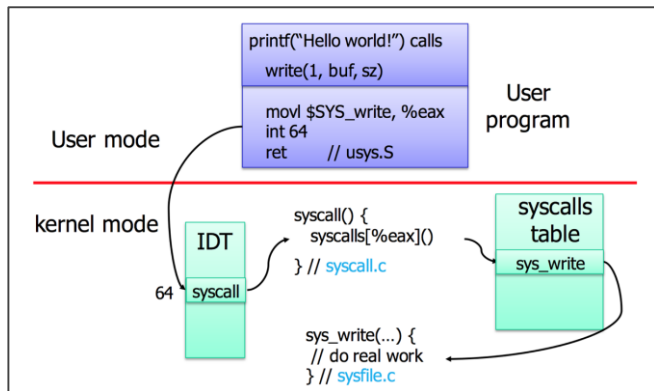
Agenda

- Brief intro to syscalls/filters/seccomp
- Overview of seccomp filter bypasses (increasing in complexity)

Background

What are Syscalls?

- The main interface between userland processes and functionality in the kernel
- The functions you call in C programs like `open`, `read`, and `write` are libc wrappers around syscalls of the same name
- Syscalls are identified by number (which changes based on the architecture)



Useful way of thinking about the userland-syscall-kernel boundary ([source](#))

What are Syscall Filters?

- There's a few ways a syscall filter can be implemented
 - Historically, some people rolled their own via [prctl](#)
 - Great talk: [Escaping the \(sand\)box from CONFidence 2017](#)
- State of the art: `seccomp-bpf` and [libseccomp](#)
 - See [here](#) and [here](#) for relevant LWN articles
 - Compiled BPF program that executes in the kernel each time a syscall is requested
 - Allows for prohibiting syscalls and their argument values
- How do we bypass these?
 - Find flaws in the implementation of the syscall filter
 - Find flaws in the set of permissible syscalls

```
A = arch
A == ARCH_I386 ? next : kill

A = sys_number
A == fork ? kill : next
A == execve ? kill : next
// ... snip ...
A >= clone ? kill : next

A == open ? check_perm : addr_check

check_perm:
A = args[2]
A == 0 ? addr_check : kill

addr_check:
A == read ? check_arg : next
A == readdir ? check_arg : next
A == write ? check_arg : allow

check_arg:
A = args[1]
A < 0x30000000 ? kill : allow

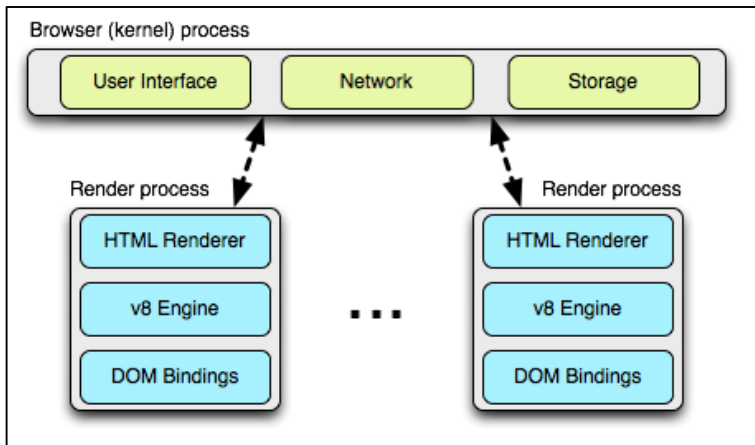
allow:
return ALLOW

kill:
return KILL_PROCESS
```

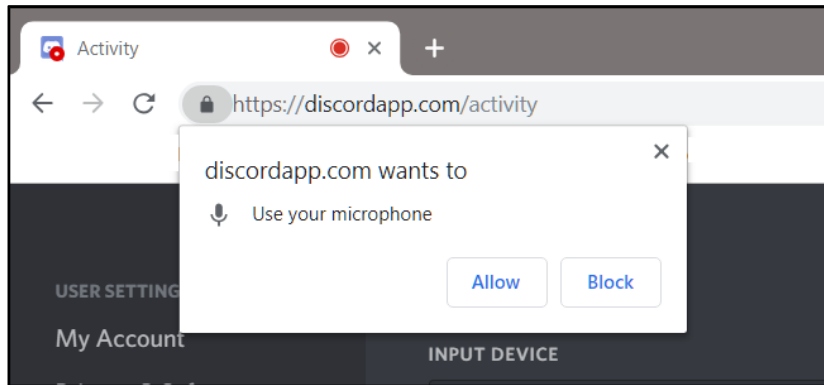
Example of a syscall filter implemented in the [seccomp-tools](#) metalanguage for a 32-bit x86 program

Defense in Depth: Syscall Filtering

- Defending against adversaries that already have code execution
- Overall goal is to limit the kernel attack surface
- Browsers are useful case study for syscall filtering
 - Helpful resource: [Chromium's Linux sandbox design document](#)



Simplified view of Chrome's multi-process architecture ([source](#))

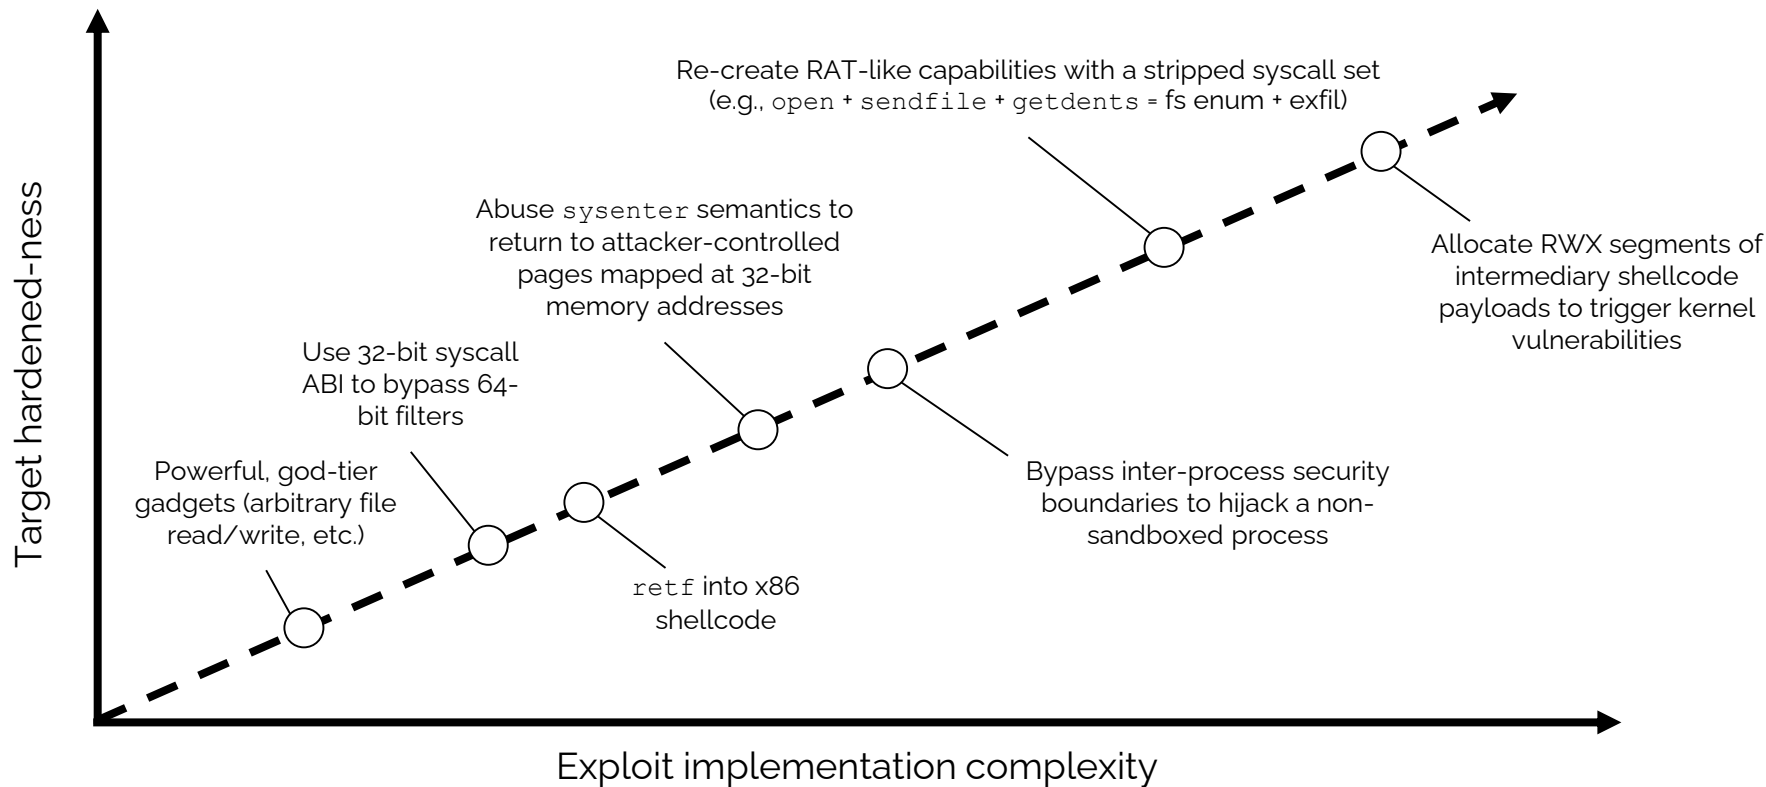


Example of granular access to system resources in the Chromium security architecture

Why Do We Need to Understand This?

- Syscall filters are the new CTF meta
 - Overwriting [free hook](#) with a one-gadget is old news
- Systems are becoming increasingly hardened; we need to stay ahead of the curve
 - Point and click access opportunities without advanced protections are a dying breed
 - Limited-execution environments guide us towards system exploitation rather than single-process exploitation

The Spectrum of Syscall Filter “Bypasses”



(Relatively) Easy Workarounds

Powerful Syscalls and Easy Wins

- Some syscalls (`fork`, `execve`, etc.) have the side effect of (sometimes) completely removing the current process's syscall filter
 - This only works for poorly-implemented `ptrace`-based filters; you're mostly safe if you use `libseccomp`
- Fortunately, some sane protections exist:
 - `seccomp-bpf` programs loaded in the kernel are preserved across `execve` and `fork` calls
- Other very dangerous syscalls:
 - Arbitrary arguments to `open/read/write`: File read/write, [procfs](#) shenanigans, etc.
 - [fork/ptrace](#): Intercept and manipulate syscalls on their way in to the kernel (example [here](#))
 - [process vm readv](#) and [process vm writev](#): Access to other processes' address spaces

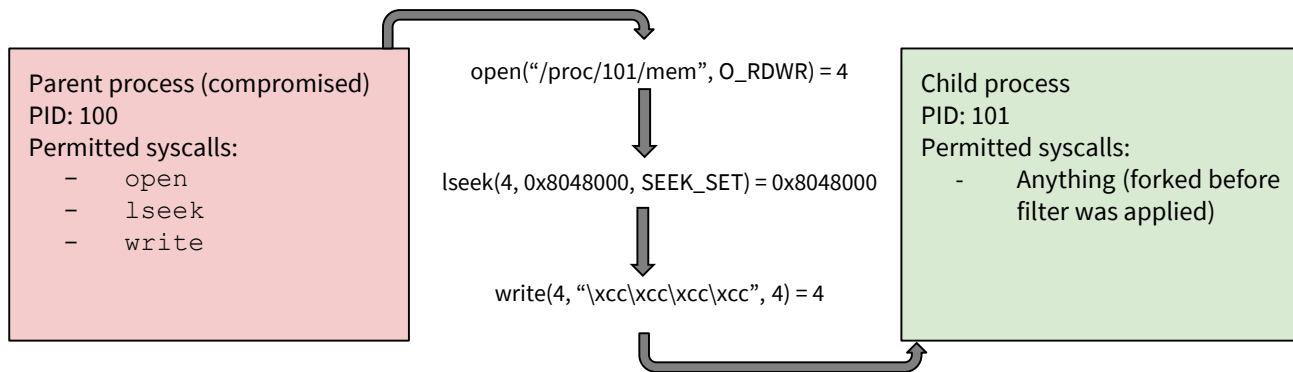
32-bit ABI Hacks

- When executing 64-bit x86 instructions, you can still execute syscalls using the 32-bit syscall ABI (i.e., via `int 0x80`)
 - These syscall numbers are *different* than their 64-bit ABI equivalents (bitwise ORed with `0x40000000`)
 - Remember that seccomp filters are implemented by checking for syscall number equivalence
- Similar (but different) technique involves switching the processor into 32-bit mode
 - Syscall numbers in 32-bit x86 mode are completely different than their 64-bit equivalents
 - Do this by executing a [far return \(`retf`\) instruction](#) with the stored CS register set for 32-bit mode
 - This can also be triggered when returning from a `sysenter` instruction issued in 64-bit mode
- This is why you shouldn't implement your own filters
 - `libseccomp` (generally) takes care of these architecture edge-case details

Entering the Arcane

Overwrite Code with `/proc/{PID}/mem`

- Exploit a sandboxed parent process → modify an un-sandboxed child process
 - Does require the ability to [lseek](#) around the opened pseudofile
- Can also modify your own process with `/proc/self/mem`
 - Not necessarily as useful, since you're still beholden to the same syscall filter
- Writing to `/proc/{PID}/mem` ignores the mapped permissions!
 - Relevant CTF challenge: [writeonly from GoogleCTF 2020](#)



Making the Most of What You Have

- Sometimes, a highly restrictive filter is still sufficient to do what we want
 - Do we really need *arbitrary* code execution?
- Minimal file system enumeration and file retrieval:
 - Open files via [open](#) or [openat](#)
 - Enumerate directory entries from `open`-ed director file descriptors with [readdir](#) or [getdents](#)
 - Send files directly to a socket via [sendfile](#)
 - Assume that we can exfiltrate data using the socket we remotely exploited over
- Relevant CTF challenge from pbctf 2020 can be found [here](#)

RWX Staging for Follow-on Kernel Exploits

The Easy Ways: mmap and mprotect

- [mmap](#): Syscall to create new memory mappings
- [mprotect](#): Syscall to change the protections of existing pages of memory
- Most sane syscall blocklists would disallow these
 - Or, more specifically, disallow the PROT_READ | PROT_WRITE | PROT_EXEC argument

```
#include <string.h>
#include <sys/mman.h>

const char shellcode[] = "\xcc\xcc\xcc\xcc";

int main(int argc, char **argv)
{
    void *mem = mmap(0, sizeof(shellcode), PROT_READ|PROT_WRITE, MAP_SHARED|MAP_ANONYMOUS, -1, 0);

    memcpy(mem, shellcode, sizeof(shellcode));

    mprotect(mem, sizeof(shellcode), PROT_READ|PROT_WRITE|PROT_EXEC);

    int (*func)();
    func = (int (*)())mem;
    (int)(*func)();

    munmap(mem, sizeof(shellcode));

    return 0;
}
```

Example of a C harness program for changing the permissions of a page of memory to RWX and executing x86 shellcode stored there (adapted from [here](#))

SYS_ipc and Friends

- Set of syscalls for managing inter-process shared memory segments
 - On 32-bit x86, `SYS_ipc` is the “gateway” syscall
- Can be used to allocate RWX pages at arbitrary addresses without the use of `mmap/mprotect`
- Easily setup in a ROP chain with 2 syscalls:
 - Create a shared memory region via `SHMGET` operation with `IPC_CREATE | 0o777` flag
 - Use the returned `key_t` to load the region into the current process, using the `SHM_EXEC` flag to mark as executable
- Example usage [here](#)



Screenshot from inside your computer when you map RWX memory via `SYS_ipc`

Side Note: You Don't Always Have to BYORWX

- Scripting languages (especially jitted ones) love to allocate RWX segments
- Up until recently, WebAssembly programs in [Blink](#) and [WebKit](#) were loaded into RWX segments

```
/* allocate a memory block */
#ifdef MS_WIN32
    item = (ITEM *)VirtualAlloc(NULL,
                                count * sizeof(ITEM),
                                MEM_COMMIT,
                                PAGE_EXECUTE_READWRITE);
    if (item == NULL)
        return;
#else
    item = (ITEM *)mmap(NULL,
                        count * sizeof(ITEM),
                        PROT_READ | PROT_WRITE | PROT_EXEC,
                        MAP_PRIVATE | MAP_ANONYMOUS,
                        -1,
                        0);
    if (item == (void *)MAP_FAILED)
        return;
#endif
```

Code from the CPython ctypes module ([source](#))



Wrapping Up

Lessons Learned

- Use syscall allowlists instead of blocklists
- Use hardened libraries (i.e., `libseccomp`)
- Syscall filtering: should only be one wall in your layered defense

Future Learning Resources

- Read Linux kernel source code

Thanks for your time

Reach out for slides and/or code samples